

Full-stack Machine Learning

Alexander D'hoore

Full-stack Machine Learning

- Introduction to MLOps
- Comparison to DevOps
- Experiment Tracking
- **Lab 1**
- Data Engineering
- **Lab 2**
- Model Deployment
 - => **you are here**
 - Cloud, edge
 - Prototyping
- **Lab 3**
- Automation
 - Agents, pipelines
 - Hyperparameters
- **Lab 4**

Model Deployment

Introduction to Model Deployment

- Model deployment refers to the process of integrating a machine learning model into an existing **production environment** to make predictions based on new data.
- It ensures that the model is accessible to users or other systems for **inference**.
- Deployment is a critical step to realize the **practical value** of a machine learning model.

Importance of Deployment

- **Bridges the gap** between model development and real-world application.
- Enables **automation and scalability** of model predictions.
- Facilitates continuous **improvement and updates** of models based on new data.

Types of Deployment: Cloud vs Edge

- **Cloud Deployment:** Models are deployed on **centralized servers** accessible via the internet.
- **Edge Deployment:** Models are deployed on **local devices**, closer to where the data is generated.

Cloud Deployment

Cloud Deployment: Benefits

- Cloud deployment involves hosting models on cloud servers, allowing for remote access via **APIs**.
- **Benefits:**
 - **Scalability:** Easily handle varying loads by scaling resources up or down.
 - **Accessibility:** Models can be accessed from anywhere via the internet.
 - **Maintenance:** Simplifies updates and management of models.
 - **Cost Efficiency:** Pay-as-you-go pricing models and resource sharing.

Cloud Deployment: Use Cases

- **E-commerce:** Product recommendation systems.
- **Healthcare:** Predictive analytics for patient data.
- **Finance:** Fraud detection and risk assessment.
- **Retail:** Demand forecasting and inventory management.

REST APIs in Cloud Deployment

- **What are REST APIs?**
- Representational State Transfer (REST) APIs
- allow systems to communicate over **HTTP**
- using standard methods like **GET, POST, PUT, and DELETE**
- They enable interaction with machine learning models deployed on cloud servers.

Role of REST APIs

- Provide a **standardized** way to expose model prediction functionality.
- Allow different applications to request and receive **model predictions**.
- Facilitate **integration** with web and mobile applications.
- **Example Workflow:** Model as a REST API
 1. Develop and train the machine learning model.
 2. Deploy the model on a cloud server.
 3. Wrap the model inference logic in a REST API.
 4. Clients send data to the API endpoint and receive predictions in response.

Cloud Providers for Model Deployment

- Overview of Popular Cloud Providers:
 - AWS (Amazon Web Services)
 - GCP (Google Cloud Platform)
 - Azure (Microsoft Azure)
 - IBM Cloud
 - Oracle Cloud

Cloud Providers: Features

- General Features Offered by Cloud Providers:
- **Compute Resources:** Virtual machines, serverless functions, and container services.
- **Storage Solutions:** Databases, object storage, and data lakes.
- **Networking:** Load balancing, virtual private clouds, and content delivery networks.
- **Security:** Identity and access management, encryption, and compliance certifications.

Model Serving Engines

- Model **serving engines** are specialized platforms designed to deploy and manage machine learning models at **scale**.
- They streamline the process of serving models and handling requests.
- Key Features to Look For:
 - **Scalability**: Ability to handle high request volumes.
 - **Latency**: Low response times for real-time predictions.
 - **Monitoring**: Tools to track model performance and usage.
 - **Integration**: Compatibility with machine learning frameworks.
 - **Management**: Version control, rollbacks, and updates.

Model Serving Engines: Examples

- Examples of Model Serving Engines
 - TensorFlow Serving
 - TorchServe
 - KFServing (KServe)
 - BentoML
 - Seldon Core

Case Study: Cloud Deployment

- A supermarket wants to deploy a demand forecasting model to optimize inventory management.
- Deployment Process
 - **Data Preparation:** Clean and preprocess the historical sales data.
 - **Model Training:** Develop and train the demand forecasting model.
 - **Model Deployment:** Deploy the model to a cloud server using a serving engine.
 - **API Setup:** Create a REST API endpoint for the model.
 - **Monitoring:** Implement monitoring and logging to track model performance.

Edge Deployment

Introduction to Edge Deployment

- Edge deployment refers to running machine learning models on **local devices** rather than in the cloud.
- Benefits:
 - **Low Latency**: Faster inference times by processing data locally.
 - **Reduced Bandwidth**: Less data transmitted over the network.
 - **Privacy and Security**: Sensitive data stays on local devices.
 - **Offline Capability**: Models can function without internet connectivity.

Possible Edge Devices

- **Mobile Devices:** Smartphones and tablets.
- **IoT Devices:** Sensors, smart home devices, and wearables.
- **Embedded Systems:** Industrial controllers, automotive systems.
- **Edge Servers:** Local servers with limited compute resources.
- **Single-Board Computers:** Raspberry Pi, NVIDIA Jetson.

TensorFlow Lite

- TensorFlow Lite is a **lightweight solution** for deploying TensorFlow models on mobile and embedded devices.
- **Key Features**
 - **Model Conversion:** Convert TensorFlow models to a compact format.
 - **Interpreter:** Execute models on-device with minimal footprint.
 - **Delegates:** Hardware acceleration using GPU, DSP, or other specialized processors.
- **Use Cases**
 - Real-time image classification on mobile phones.
 - Speech recognition in embedded systems.
 - Object detection on IoT devices.

ONNX (Open Neural Network Exchange) Format

- ONNX is an **open standard** for representing machine learning models.
 - Enables interoperability between different frameworks.
- **Key Features**
 - **Model Interoperability:** Transfer models between frameworks like PyTorch, TensorFlow, and others.
 - **Optimization:** Tools to optimize models for performance on various hardware.
- **Use Cases**
 - Converting models developed in PyTorch to run on different platforms.
 - Sharing models across teams using different machine learning frameworks.

Microsoft ONNX Runtime

- ONNX Runtime is a cross-platform, **high-performance inference engine**.
 - For Open Neural Network Exchange (ONNX) models.
- **Key Features**
 - **Cross-Platform:** Runs on Windows, Linux, Mac, and various edge devices.
 - **Optimizations:** Hardware-specific optimizations for CPUs, GPUs, and other accelerators.
 - **Extensibility:** Custom operators and execution providers.
- **Use Cases**
 - Deploying complex models on powerful edge device.
 - Integrating ONNX models into existing applications for real-time inference.

NVIDIA TensorRT

- TensorRT is a **high-performance** deep learning inference optimizer and runtime library developed by NVIDIA.
- **Key Features**
 - **Model Optimization:** Converts and optimizes models for NVIDIA GPUs.
 - **Low Latency:** Minimizes inference time with advanced optimizations.
 - **Precision Calibration:** Supports FP16 and INT8 precision for faster performance.
- **Use Cases**
 - Deploying deep learning models on NVIDIA Jetson for robotics.
 - Accelerating AI applications in autonomous vehicles.
 - Real-time video analytics on edge devices.

Best Practices for Edge Deployment

- Model Optimization
 - **Quantization:** Reduce model size by converting weights to lower precision (e.g., INT8).
 - **Pruning:** Remove redundant neurons and weights.
 - **Compression:** Use techniques like weight clustering and Huffman coding.
- Resource Management
 - **Memory:** Optimize memory usage to fit models on limited hardware.
 - **Compute:** Balance model complexity with available processing power.
 - **Energy:** Minimize energy consumption for battery-powered devices.

Case Study: Edge Deployment

- A smart home company wants to deploy a voice recognition model on their IoT devices to enable voice commands.
- Deployment Process
 - **Data Preparation:** Collect and preprocess audio data for training.
 - **Model Training:** Develop and train the voice recognition model.
 - **Model Conversion:** Convert the model to TensorFlow Lite format.
 - **Device Deployment:** Deploy the model on IoT devices.
 - **Integration:** Integrate the model with the device firmware to handle voice commands.

Summary: Cloud vs Edge Deployment

- Cloud Deployment
 - **Scalability:** Easily handles varying loads
 - **Accessibility:** Global access via the internet
 - **Maintenance:** Simplified updates
 - **Challenges:** Higher latency, bandwidth costs, data privacy
- Edge Deployment
 - **Low Latency:** Fast local processing
 - **Privacy:** Data stays on device
 - **Offline Capability:** Works without internet
 - **Challenges:** Limited resources, harder maintenance

Prototyping Apps

Introduction to Prototyping Apps

- Prototyping apps allow **rapid development** and deployment of interactive machine learning demos and applications.
- Use Cases:
 - **Internal Demos:** Quickly create interactive apps to showcase models and data insights.
 - **Custom Labeling:** Develop labeling tools for data annotation tasks.
 - **Monitoring Dashboards:** Build comprehensive dashboards to track metrics and visualize data.

Prototyping Apps: Examples

- **Image Classification:** Upload an image and get predictions.
- **Text Analysis:** Input text and receive sentiment analysis results.
- **Audio Processing:** Record or upload audio and get transcription.
- **Data Visualization:** Create interactive plots and charts.
- **Model Exploration:** Adjust model parameters and observe changes in output.
- **Data Analysis:** Perform exploratory data analysis with interactive controls.

Prototyping: Gradio

- Gradio is an open-source library for creating **interactive user interfaces** for machine learning models.
- Simplifies sharing models with non-technical stakeholders.
- Key Features:
 - **Easy Integration:** Quickly wrap models and functions with a web interface.
 - **Interactive Demos:** Create live demos for model inference.
 - **Customizable Interfaces:** Build interfaces with buttons, sliders, text inputs, and more.
 - **Deployment Options:** Host locally or deploy on the web with minimal effort.

Prototyping: Streamlit

- Streamlit is an open-source app framework for creating and sharing **data applications**.
- Focuses on simplicity and speed for building **custom web applications**.
- Key Features:
 - **Python-Based**: Build apps using pure Python code.
 - **Real-Time Updates**: Automatically update the app as code changes.
 - **Widgets and Controls**: Use sliders, dropdowns, and other UI elements to interact with data.
 - **Deployment Options**: Deploy locally or use Streamlit Cloud for online hosting.

Comparison: Gradio vs Streamlit

- **Gradio**

- Best for quick, interactive model demos.
- Simple to create interfaces for machine learning models.
- Ideal for non-technical stakeholders to interact with models.

- **Streamlit**

- Best for building comprehensive data applications and dashboards.
- Greater flexibility and control over the app layout and functionality.
- Suitable for interactive data visualization and analysis.

Lab 3: Deployment

Lab 3: Deployment

- We will learn how to build **Gradio apps**
- Follow the tutorial at:
<https://www.gradio.app/guides/quickstart>